


```
# STEP 1: Install & authenticate Kaggle
!pip install -q kaggle

# Upload your Kaggle API key (kaggle.json)
from google.colab import files
files.upload() # Upload the kaggle.json file

# Create kaggle folder and move the key
!mkdir -p ~/.kaggle
!cp kaggle.json ~/.kaggle/
!chmod 600 ~/.kaggle/kaggle.json

# STEP 2: Download the Animals-10 dataset
!kaggle datasets download -d alessiocrrado99/animals10

# Unzip the dataset
!unzip -q animals10.zip -d animals10
```

 No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving kaggle.json to kaggle.json
 Dataset URL: <https://www.kaggle.com/datasets/alessiocrrado99/animals10>
 License(s): GDF-2.0

```
translate = {
    "cane": "dog",
    "cavallo": "horse",
    "elefante": "elephant",
    "farfalla": "butterfly",
    "gallina": "chicken",
    "gatto": "cat",
    "mucca": "cow",
    "pecora": "sheep",
    "ragno": "spider",
    "scoiattolo": "squirrel",
}
```


```
directory=r"/content/animals10/raw-img"
batch_size = 32
img_size = 224
```

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator, load_img, img_to_array
```

```
datagen = ImageDataGenerator(
    rescale=1/255.,
    horizontal_flip=True,
    validation_split=0.15
)
```

```
train_generator = datagen.flow_from_directory(
    directory,
    target_size=(img_size, img_size),
    batch_size=batch_size,
    shuffle=True,
    subset='training',
    class_mode='categorical'
)
```

```
validation_generator = datagen.flow_from_directory( directory,
    target_size=(img_size, img_size),
    batch_size=batch_size,
    shuffle=False,
    subset='validation',
    class_mode='categorical'
)
```

 Found 22257 images belonging to 10 classes.
 Found 3922 images belonging to 10 classes.

```
import numpy as np
import matplotlib.pyplot as plt
```

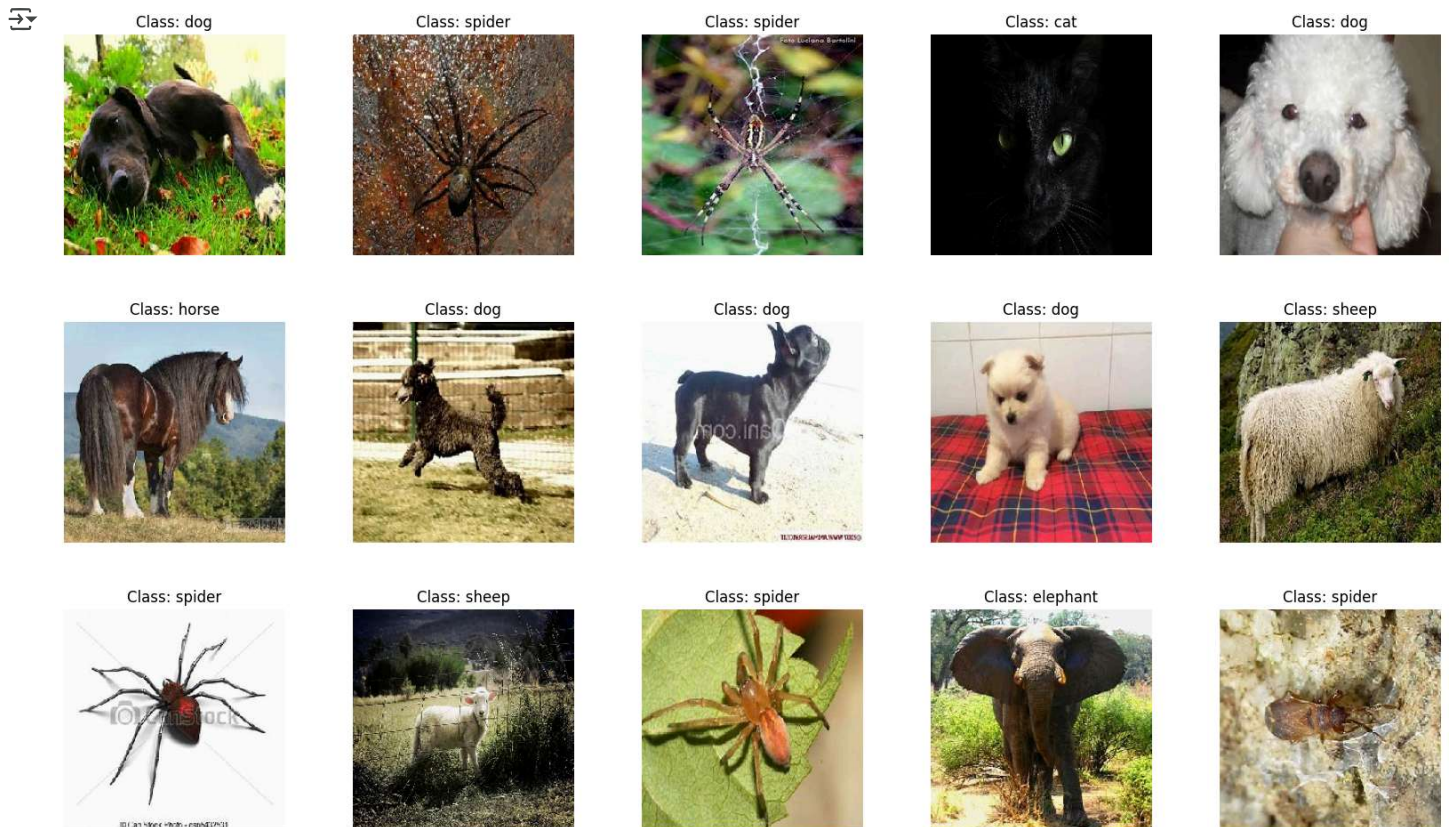
```

labels = [k for k in train_generator.class_indices]
translated_labels = [translate[label] for label in labels]
sample_generate = train_generator.__next__()
images = sample_generate[0]
titles = sample_generate[1]
plt.figure(figsize=(20, 20))

for i in range(15):
    plt.subplot(5, 5, i+1)
    plt.subplots_adjust(hspace=0.3, wspace=0.3)
    plt.imshow(images[i])
    plt.title(f'Class: {translated_labels[np.argmax(titles[i], axis=0)]}')
    plt.axis("off")

plt.show()

```



```

import tensorflow as tf
from tensorflow.keras.layers import (Conv2D, MaxPooling2D, AveragePooling2D, Flatten, Dense, Input, Dropout,
                                     concatenate)
from tensorflow.keras.models import Sequential, Model
from tensorflow.keras.applications import VGG16
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.utils import to_categorical
import matplotlib.pyplot as plt
import numpy as np
import os
import time

img_size = 224

base_model = VGG16(include_top=False,
                    weights='imagenet',

```

```
input_shape=(img_size, img_size, 3))
```

```
base_model.summary()
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.58889256/58889256 0s 0us/step
Model: "vgg16"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1,792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36,928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73,856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147,584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295,168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590,080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590,080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1,180,160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0

Total params: 14,714,688 (56.13 MB)

Trainable params: 14,714,688 (56.13 MB)

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
```

```
def build_zfnet():
    model = Sequential([
        Conv2D(96, (7,7), strides=(2,2), activation='relu', input_shape=(224,224,3)),
        MaxPooling2D((3,3), strides=(2,2)),
        Conv2D(256, (5,5), strides=(2,2), activation='relu'),
        MaxPooling2D((3,3), strides=(2,2)),
        Conv2D(384, (3,3), activation='relu'),
        Conv2D(384, (3,3), activation='relu'),
        Conv2D(256, (3,3), activation='relu'),
        MaxPooling2D((3,3), strides=(2,2)),
        Flatten(),
        Dense(4096, activation='relu'),
        Dense(4096, activation='relu'),
        Dense(10, activation='softmax')
    ])
    model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
    return model
```

```
from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Model
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense
```

```
def build_vgg16():
    base = VGG16(weights='imagenet', include_top=False, input_shape=(224,224,3))
    for layer in base.layers:
        layer.trainable = False
```

```

x = GlobalAveragePooling2D()(base.output)
x = Dense(256, activation='relu')(x)
output = Dense(10, activation='softmax')(x)
model = Model(base.input, output)
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
return model

import tensorflow as tf
from tensorflow.keras.layers import Conv2D, MaxPooling2D, AveragePooling2D, Flatten, Dense, Input, concatenate
from tensorflow.keras.models import Model
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import matplotlib.pyplot as plt

import tensorflow as tf
from tensorflow.keras.layers import Conv2D, MaxPooling2D, AveragePooling2D, Flatten, Dense, Input, concatenate
from tensorflow.keras.models import Model
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import matplotlib.pyplot as plt

# --- Define Inception Module ---
def inception_module(x, f1, f3r, f3, f5r, f5, pool_proj):
    conv1 = Conv2D(f1, (1,1), padding='same', activation='relu')(x)
    conv3 = Conv2D(f3r, (1,1), padding='same', activation='relu')(x)
    conv3 = Conv2D(f3, (3,3), padding='same', activation='relu')(conv3)
    conv5 = Conv2D(f5r, (1,1), padding='same', activation='relu')(x)
    conv5 = Conv2D(f5, (5,5), padding='same', activation='relu')(conv5)
    pool = MaxPooling2D((3,3), strides=(1,1), padding='same')(x)
    pool = Conv2D(pool_proj, (1,1), padding='same', activation='relu')(pool)
    output = concatenate([conv1, conv3, conv5, pool], axis=3)
    return output

# --- Build Inception V1 Model ---
def build_inception_v1(input_shape=(224, 224, 3), num_classes=10):
    input_layer = Input(shape=input_shape)

    x = Conv2D(64, (7,7), strides=(2,2), padding='same', activation='relu')(input_layer)
    x = MaxPooling2D((3,3), strides=(2,2), padding='same')(x)
    x = Conv2D(64, (1,1), activation='relu')(x)
    x = Conv2D(192, (3,3), padding='same', activation='relu')(x)
    x = MaxPooling2D((3,3), strides=(2,2), padding='same')(x)

    x = inception_module(x, 64, 96, 128, 16, 32, 32)
    x = inception_module(x, 128, 128, 192, 32, 96, 64)
    x = MaxPooling2D((3,3), strides=(2,2), padding='same')(x)

    x = inception_module(x, 192, 96, 208, 16, 48, 64)
    x = inception_module(x, 160, 112, 224, 24, 64, 64)
    x = inception_module(x, 128, 128, 256, 24, 64, 64)
    x = inception_module(x, 112, 144, 288, 32, 64, 64)
    x = inception_module(x, 256, 160, 320, 32, 128, 128)
    x = MaxPooling2D((3,3), strides=(2,2), padding='same')(x)

    x = inception_module(x, 256, 160, 320, 32, 128, 128)
    x = inception_module(x, 384, 192, 384, 48, 128, 128)

    x = AveragePooling2D(pool_size=(7,7), strides=(1,1))(x)
    x = Flatten()(x)
    x = Dense(1024, activation='relu')(x)
    x = Dense(num_classes, activation='softmax')(x)

    model = Model(input_layer, x)
    model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
    return model

models = {
    "ZFNet": build_zfnet(),
    "VGG16": build_vgg16(),
    "InceptionV1": build_inception_v1()
}

/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape` / `input_dim` argument to `Layer.__init__` (it is handled by `Layer.make_node` instead).
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```

```

import time
import matplotlib.pyplot as plt
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
histories = {}
results = {}

for name, model in models.items():
    print(f"\n--- Training {name} ---")
    start_time = time.time()
    default_epochs = 20
    callbacks = []

    if name.lower() == "zfnet":
        epochs = 20
        callbacks = [
            EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True),
            ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=2, verbose=1)
        ]
    else:
        epochs = default_epochs
        callbacks = [
            EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)
        ]

    history = model.fit(
        train_generator,
        validation_data=validation_generator,
        epochs=epochs,
        verbose=1,
        callbacks=callbacks
    )

    end_time = time.time()

    val_loss, val_acc = model.evaluate(validation_generator, verbose=0)
    print(f"{name} Validation Accuracy: {val_acc:.4f}")
    print(f"{name} Training Time: {end_time - start_time:.2f} seconds")

    histories[name] = history
    results[name] = {
        'val_accuracy': val_acc,
        'training_time': end_time - start_time
    }

    save_path = f'{name}_model.keras'
    model.save(save_path)
    print(f"{name} model saved to {save_path}")

def plot_histories(histories):
    for name, history in histories.items():
        plt.figure(figsize=(12, 5))
        plt.suptitle(f'{name} Training Performance', fontsize=16)

        # Accuracy
        plt.subplot(1, 2, 1)
        plt.plot(history.history['accuracy'], label='Train Accuracy')
        plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
        plt.title('Accuracy')
        plt.xlabel('Epochs')
        plt.ylabel('Accuracy')
        plt.legend()

        # Loss
        plt.subplot(1, 2, 2)
        plt.plot(history.history['loss'], label='Train Loss')
        plt.plot(history.history['val_loss'], label='Validation Loss')
        plt.title('Loss')
        plt.xlabel('Epochs')
        plt.ylabel('Loss')
        plt.legend()

        plt.show()

plot_histories(histories)

```

```
print("\n--- Final Evaluation Summary ---")
for name, data in results.items():
    print(f"{name}:")
    print(f"  Validation Accuracy: {data['val_accuracy']:.4f}")
    print(f"  Training Time: {data['training_time']:.2f} seconds\n")
```



```
--- Training ZFNet ---
/usr/local/lib/python3.11/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class
  self._warn_if_super_not_called()
Epoch 1/20
696/696 ██████████ 74s 94ms/step - accuracy: 0.2261 - loss: 2.2471 - val_accuracy: 0.3343 - val_loss: 1.8556 - learning_rate:
Epoch 2/20
696/696 ██████████ 62s 90ms/step - accuracy: 0.3550 - loss: 1.8385 - val_accuracy: 0.3921 - val_loss: 1.7264 - learning_rate:
Epoch 3/20
696/696 ██████████ 59s 84ms/step - accuracy: 0.4067 - loss: 1.6904 - val_accuracy: 0.4151 - val_loss: 1.6455 - learning_rate:
Epoch 4/20
696/696 ██████████ 61s 88ms/step - accuracy: 0.4655 - loss: 1.5416 - val_accuracy: 0.4906 - val_loss: 1.4824 - learning_rate:
Epoch 5/20
696/696 ██████████ 61s 88ms/step - accuracy: 0.5242 - loss: 1.4050 - val_accuracy: 0.5709 - val_loss: 1.2619 - learning_rate:
Epoch 6/20
696/696 ██████████ 57s 82ms/step - accuracy: 0.5604 - loss: 1.2975 - val_accuracy: 0.5745 - val_loss: 1.2563 - learning_rate:
Epoch 7/20
696/696 ██████████ 59s 85ms/step - accuracy: 0.5822 - loss: 1.2384 - val_accuracy: 0.6045 - val_loss: 1.1593 - learning_rate:
Epoch 8/20
696/696 ██████████ 58s 83ms/step - accuracy: 0.6002 - loss: 1.1827 - val_accuracy: 0.6114 - val_loss: 1.1510 - learning_rate:
Epoch 9/20
696/696 ██████████ 62s 89ms/step - accuracy: 0.6246 - loss: 1.1060 - val_accuracy: 0.6334 - val_loss: 1.0931 - learning_rate:
Epoch 10/20
696/696 ██████████ 60s 87ms/step - accuracy: 0.6404 - loss: 1.0590 - val_accuracy: 0.6397 - val_loss: 1.0519 - learning_rate:
Epoch 11/20
696/696 ██████████ 58s 83ms/step - accuracy: 0.6619 - loss: 1.0062 - val_accuracy: 0.6543 - val_loss: 1.0477 - learning_rate:
Epoch 12/20
696/696 ██████████ 58s 83ms/step - accuracy: 0.6725 - loss: 0.9591 - val_accuracy: 0.6548 - val_loss: 1.0339 - learning_rate:
Epoch 13/20
696/696 ██████████ 59s 85ms/step - accuracy: 0.6917 - loss: 0.9141 - val_accuracy: 0.6586 - val_loss: 1.0300 - learning_rate:
Epoch 14/20
696/696 ██████████ 58s 83ms/step - accuracy: 0.7044 - loss: 0.8771 - val_accuracy: 0.6708 - val_loss: 1.0239 - learning_rate:
Epoch 15/20
696/696 ██████████ 57s 82ms/step - accuracy: 0.7134 - loss: 0.8535 - val_accuracy: 0.6711 - val_loss: 1.0146 - learning_rate:
Epoch 16/20
696/696 ██████████ 84s 84ms/step - accuracy: 0.7209 - loss: 0.8344 - val_accuracy: 0.6795 - val_loss: 0.9792 - learning_rate:
Epoch 17/20
696/696 ██████████ 62s 88ms/step - accuracy: 0.7330 - loss: 0.7855 - val_accuracy: 0.6902 - val_loss: 0.9513 - learning_rate:
Epoch 18/20
696/696 ██████████ 59s 85ms/step - accuracy: 0.7495 - loss: 0.7424 - val_accuracy: 0.6527 - val_loss: 1.1038 - learning_rate:
Epoch 19/20
696/696 ██████████ 0s 72ms/step - accuracy: 0.7518 - loss: 0.7374
Epoch 19: ReduceLROnPlateau reducing learning rate to 0.0005000000237487257.
696/696 ██████████ 56s 81ms/step - accuracy: 0.7518 - loss: 0.7374 - val_accuracy: 0.6887 - val_loss: 0.9757 - learning_rate:
Epoch 20/20
696/696 ██████████ 58s 83ms/step - accuracy: 0.7968 - loss: 0.5949 - val_accuracy: 0.7104 - val_loss: 0.9704 - learning_rate:
ZFNet Validation Accuracy: 0.6920
ZFNet Training Time: 1223.46 seconds
ZFNet model saved to ZFNet_model.keras
```

```
--- Training VGG16 ---
Epoch 1/20
696/696 ██████████ 183s 242ms/step - accuracy: 0.5806 - loss: 1.2896 - val_accuracy: 0.7772 - val_loss: 0.6776
Epoch 2/20
696/696 ██████████ 149s 214ms/step - accuracy: 0.7932 - loss: 0.6285 - val_accuracy: 0.8027 - val_loss: 0.5926
Epoch 3/20
696/696 ██████████ 149s 214ms/step - accuracy: 0.8246 - loss: 0.5380 - val_accuracy: 0.8108 - val_loss: 0.5553
```

```
from sklearn.metrics import confusion_matrix, classification_report, ConfusionMatrixDisplay
import numpy as np
import matplotlib.pyplot as plt

def evaluate_model(model, generator, class_names, model_name):
```

```
    true_labels = generator.classes
    predictions = model.predict(generator, verbose=1)
    predicted_labels = np.argmax(predictions, axis=1)
    cm = confusion_matrix(true_labels, predicted_labels)
    plt.figure(figsize=(10, 8))
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=class_names)
    disp.plot(cmap='Blues', xticks_rotation='vertical', values_format='d')
    plt.title(f'{model_name} Confusion Matrix')
    plt.show()
    report = classification_report(true_labels, predicted_labels, target_names=class_names)
```



```

print(f"--- {model_name} Classification Report ---\n")
print(report)
class_names = list(validation_generator.class_indices.keys())


for name, model in models.items():
    print(f"\n--- Evaluating {name} ---")
    evaluate_model(model, validation_generator, class_names, name)

```

```

Epoch 18/20
696/696 --- 149s 214ms/step - accuracy: 0.8959 - loss: 0.3053 - val_accuracy: 0.8623 - val_loss: 0.4294
Epoch 19/20
696/696 --- 149s 214ms/step - accuracy: 0.8959 - loss: 0.3053 - val_accuracy: 0.8623 - val_loss: 0.4294
Epoch 20/20
696/696 --- 149s 214ms/step - accuracy: 0.8959 - loss: 0.3053 - val_accuracy: 0.8623 - val_loss: 0.4294
VGG16 Confusion Matrix

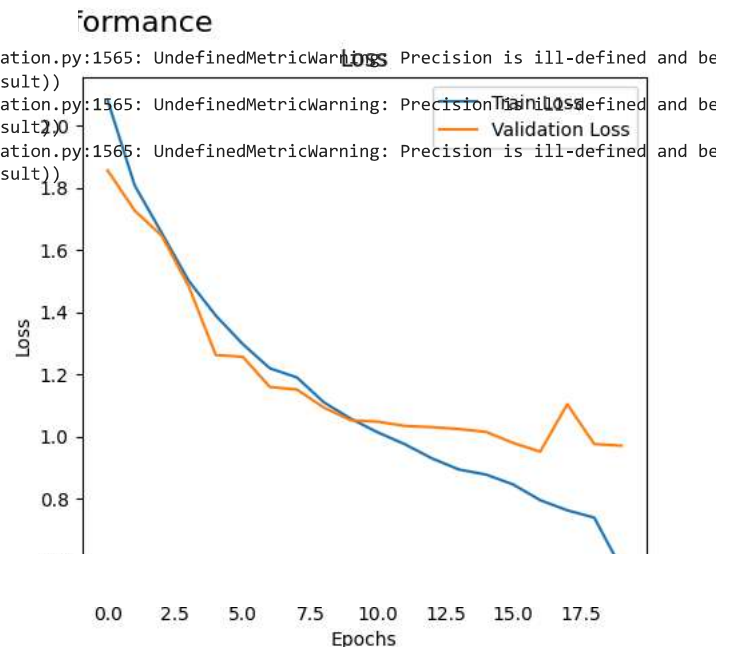
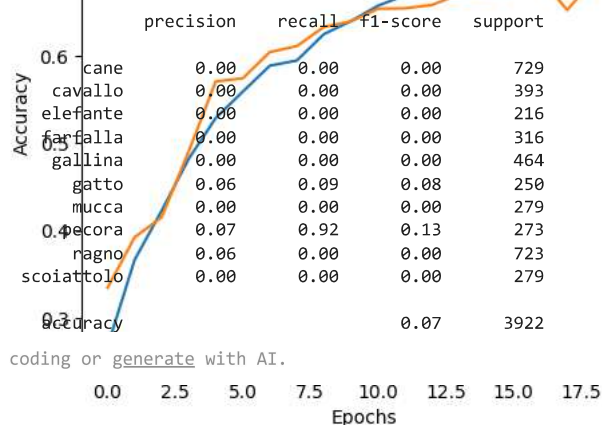
```



```

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and be
_warn_pre(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and be
_warn_pre(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and be
_warn_pre(average, modifier, f"{metric.capitalize()} is", len(result))
--- ZFNet Classification Report ---

```




Start coding or [generate](#) with AI.

```

--- Evaluating VGG16 ---
123/123 --- 2279s 19s/step
<Figure size 1000x800 with 0 Axes>

```

VGG16 Training Performance



Loss

— Train Loss
— Validation Loss

